

```
In [42]: import numpy as np
import tensorflow as tf
from tensorflow import keras
import pandas as pd
from matplotlib import pyplot as plt
%matplotlib inline
import os
import cv2
from keras.layers import Dense, Flatten
```

```
In [2]: # Define the paths to your image folders
train_path = 'C:\\Users\\biraj\\OneDrive\\PythonProjects\\Machine learning neural network\\'
val_path = 'C:\\Users\\biraj\\OneDrive\\PythonProjects\\Machine learning neural network\\'
```

```
In [3]: # Set the path to the folder containing the 'train' folder
data_dir = train_path

# Set the image size
img_size = (32, 32)

# Create empty lists for the images and labels
images = []
labels = []

# Loop over each folder from '0' to '9'
for label in range(10):
    folder_path = os.path.join(data_dir, 'train', str(label))
    # Loop over each image in the folder
    for file in os.listdir(folder_path):
        file_path = os.path.join(folder_path, file)
        if file_path.endswith(('.tiff', '.bmp')):
            # Load the image and resize it to the desired size
            img = cv2.imread(file_path, cv2.IMREAD_GRAYSCALE)
            img = cv2.resize(img, img_size)
            # Append the image and label to the lists
            images.append(img)
            labels.append(label)

# Convert the lists to NumPy arrays
```

```
images = np.array(images)
labels = np.array(labels)

# Save the arrays in NumPy format
np.save('x_train.npy', images)
np.save('y_train.npy', labels)
```

```
In [56]: # Set the path to the folder containing the 'val' folder
data_dir_val = val_path

# Set the image size
img_size_val = (32, 32)

# Create empty lists for the images and labels
images_val = []
labels_val = []

# Loop over each folder from '0' to '9'
for label in range(10):
    folder_path = os.path.join(data_dir_val, 'val\\', str(label))

    # Loop over each image in the folder
    for file in os.listdir(folder_path):
        file_path = os.path.join(folder_path, file)
        if file_path.endswith(('.tiff', '.bmp')):
            # Load the image and resize it to the desired size
            img = cv2.imread(file_path, cv2.IMREAD_GRAYSCALE)
            img = cv2.resize(img, img_size_val)
            # Append the image and label to the lists
            images_val.append(img)
            labels_val.append(label)

# Convert the lists to NumPy arrays
images_val = np.array(images_val)
labels_val = np.array(labels_val)

# Save the arrays in NumPy format
np.save('x_test.npy', images_val)
np.save('y_test.npy', labels_val)
```

```
In [82]: # Load the dataset
x_train = np.load('x_train.npy')
y_train = np.load('y_train.npy')
x_test = np.load('x_test.npy')
y_test = np.load('y_test.npy')
```

```
In [83]: # test the images are loaded correctly
```

```
print(len(x_train))
print(len(x_test))
x_train[0].shape
x_train[0]
plt.matshow(x_train[0])
plt.matshow(x_train[999])
print(x_train.shape)
print(x_test.shape)
y_train
y_test
plt.matshow(x_test[150])
```


```
1000
```


```
178
```


```
(1000, 32, 32)
```

```
(178, 32, 32)
```

```
Out[83]: <matplotlib.image.AxesImage at 0x20ba468abb0>
```







In [84]: *## flatten the dataset i.e, change 2D to 1D (skipped this , and flattened in the model)*

```
# x_train_flat = x_train.reshape(len(x_train),32*32)  
# x_test_flat = x_test.reshape(len(x_test),32*32)  
  
# print(x_train_flat.shape)  
# print(x_test_flat.shape)  
  
# x_train_flat[0]
```

In [85]: *# creating a simple nn*
create a dense layer where every input is connected to every other output, the number of inputs are 1000, outputs are 10
activation function is sigmoid

```
model = keras.Sequential([  
    keras.layers.Flatten(),
```

```

keras.layers.Dense(10, input_shape=(1024,),activation = 'sigmoid')
])

# compile the nn
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy']
            )
# train the model
# some 10 iterations done here
model.fit(x_train, y_train, epochs= 10, validation_data=(x_test, y_test))

```

```

Epoch 1/10
32/32 [=====] - 1s 13ms/step - loss: 172.8362 - accuracy: 0.3560 - val_loss: 54.1465 - val_accuracy: 0.5899
Epoch 2/10
32/32 [=====] - 0s 6ms/step - loss: 21.8444 - accuracy: 0.7990 - val_loss: 37.7204 - val_accuracy: 0.7303
Epoch 3/10
32/32 [=====] - 0s 6ms/step - loss: 14.4273 - accuracy: 0.8600 - val_loss: 33.3826 - val_accuracy: 0.7640
Epoch 4/10
32/32 [=====] - 0s 5ms/step - loss: 8.5064 - accuracy: 0.9100 - val_loss: 21.0088 - val_accuracy: 0.8090
Epoch 5/10
32/32 [=====] - 0s 5ms/step - loss: 6.6391 - accuracy: 0.9110 - val_loss: 21.1726 - val_accuracy: 0.8090
Epoch 6/10
32/32 [=====] - 0s 5ms/step - loss: 3.2933 - accuracy: 0.9510 - val_loss: 13.3256 - val_accuracy: 0.9045
Epoch 7/10
32/32 [=====] - 0s 5ms/step - loss: 1.9723 - accuracy: 0.9710 - val_loss: 16.6435 - val_accuracy: 0.8371
Epoch 8/10
32/32 [=====] - 0s 5ms/step - loss: 1.8374 - accuracy: 0.9650 - val_loss: 16.0006 - val_accuracy: 0.8483
Epoch 9/10
32/32 [=====] - 0s 5ms/step - loss: 6.0564 - accuracy: 0.9380 - val_loss: 13.9050 - val_accuracy: 0.8933
Epoch 10/10
32/32 [=====] - 0s 5ms/step - loss: 2.2400 - accuracy: 0.9710 - val_loss: 19.2359 - val_accuracy: 0.8820

```

Out[85]: <keras.callbacks.History at 0x20ba7a2ee20>

In [86]: *# Observation : we see a better accuracy from the 2nd iteration*

In [87]: *# now scale and try to check the accuracy, divide dataset by 255*

```
x_train_scaled = x_train/255
x_test_scaled = x_test/255
```

```
model.fit(x_train_scaled, y_train, epochs=10, validation_data=(x_test_scaled, y_test))
```

Epoch 1/10

32/32 [=====] - 1s 13ms/step - loss: 0.7792 - accuracy: 0.9040 - val_loss: 1.0426 - val_accuracy: 0.7865

Epoch 2/10

32/32 [=====] - 0s 5ms/step - loss: 0.7721 - accuracy: 0.8990 - val_loss: 0.9965 - val_accuracy: 0.8202

Epoch 3/10

32/32 [=====] - 0s 5ms/step - loss: 0.7343 - accuracy: 0.9380 - val_loss: 0.9635 - val_accuracy: 0.8596

Epoch 4/10

32/32 [=====] - 0s 5ms/step - loss: 0.7079 - accuracy: 0.9630 - val_loss: 0.9408 - val_accuracy: 0.9045

Epoch 5/10

32/32 [=====] - 0s 5ms/step - loss: 0.6901 - accuracy: 0.9760 - val_loss: 0.9264 - val_accuracy: 0.9157

Epoch 6/10

32/32 [=====] - 0s 5ms/step - loss: 0.6772 - accuracy: 0.9860 - val_loss: 0.9150 - val_accuracy: 0.9213

Epoch 7/10

32/32 [=====] - 0s 5ms/step - loss: 0.6671 - accuracy: 0.9860 - val_loss: 0.9070 - val_accuracy: 0.9270

Epoch 8/10

32/32 [=====] - 0s 5ms/step - loss: 0.6584 - accuracy: 0.9910 - val_loss: 0.8993 - val_accuracy: 0.9213

Epoch 9/10

32/32 [=====] - 0s 5ms/step - loss: 0.6513 - accuracy: 0.9920 - val_loss: 0.8931 - val_accuracy: 0.9157

Epoch 10/10

32/32 [=====] - 0s 5ms/step - loss: 0.6446 - accuracy: 0.9920 - val_loss: 0.8866 - val_accuracy: 0.9213

Out[87]: <keras.callbacks.History at 0x20ba3f89ca0>

In [21]: *# Observation : we got better result for all iterations on scaling the training dataset*

In [88]: *# evaluate test dataset*

```
model.evaluate(x_test_scaled, y_test)
```

6/6 [=====] - 0s 4ms/step - loss: 0.8866 - accuracy: 0.9213

Out[88]: [0.8866074681282043, 0.9213483333587646]


```
In [23]: # Observation : result almost same as the training dataset,
```

```
In [90]: # predict 1st image
```

```
plt.matshow(x_test[0])
```

```
y_predicted = model.predict(x_test_scaled)  
y_predicted[0]
```

```
# this showing the 10 results for the input '0', we need to look for the value which is max
```

```
print('Predicted Value is ', np.argmax(y_predicted[0]))
```

```
# test some more values
```

```
plt.matshow(x_test[88])
```

```
print('Predicted Value is ', np.argmax(y_predicted[88]))
```

```
plt.matshow(x_test[177])
```


```
print('Predicted Value is ', np.argmax(y_predicted[177]))
```


```
6/6 [=====] - 0s 3ms/step
```


```
Predicted Value is 0
```

```
Predicted Value is 5
```

```
Predicted Value is 9
```







```
In [91]: # some predictions may not be not right
```

```
# build confusion matrix to see how our prediction looks like
```

```
# convert to concrete values
```

```
y_predicted_labels=[np.argmax(i) for i in y_predicted]
```

```
print(y_predicted_labels, len(y_predicted_labels))
```

```
conf_mat = tf.math.confusion_matrix(labels=y_test, predictions=y_predicted_labels)
```


```
conf_mat
```

[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 7, 1, 1, 1, 1, 1, 1, 1, 1, 7, 1, 7, 1, 1, 7, 2, 2, 7, 2, 2, 2, 2, 3,
2, 2, 2, 2, 3, 4, 7, 2, 5, 3, 3, 3, 3, 3, 3, 3, 3, 3, 9, 3, 3, 3, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
4, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 1, 7, 7, 7, 7, 7,
7, 7, 7, 7, 1, 7, 7, 7, 7, 7, 1, 7, 8, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9,
9, 9, 9, 9, 9, 9] 178

```
Out[91]: <tf.Tensor: shape=(10, 10), dtype=int32, numpy=
array([[18,  0,  0,  0,  0,  0,  0,  0,  0,  0],
       [ 0, 12,  0,  0,  0,  0,  0,  4,  0,  0],
       [ 0,  0, 11,  2,  1,  1,  0,  2,  0,  0],
       [ 0,  0,  0, 16,  0,  0,  0,  0,  0,  1],
       [ 0,  0,  0,  0, 18,  0,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0, 18,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0,  0, 18,  0,  0,  0],
       [ 0,  3,  0,  0,  0,  0,  0, 15,  0,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0, 18,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0,  0, 20]])>
```

```
In [92]: import seaborn as sn
plt.figure(figsize = (10,10))
sn.heatmap(conf_mat,annot=True,fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
Out[92]: Text(95.7222222222221, 0.5, 'Actual')
```



```
In [28]: # here we can see there are some errors  
# we need to modify our nn, we add some layers in the above model and different activation function
```

```
In [116... # in 1st Dense layer, the input is 32 x 32 = 1024 neurons, which will give 10 output (numbers from 0 to 9)  
# 2nd Dense layer, the input is 10 neurons from above layers output  
# we can add more layers for accuracy  
  
model2 = keras.Sequential([  
    keras.layers.Flatten(),  
    keras.layers.Dense(1024, input_shape=(1024,)), activation='relu'),  
    keras.layers.Dense(10, activation='softmax')  
)  
  
# compile the nn  
model2.compile(optimizer='adam',  
               loss='sparse_categorical_crossentropy',  
               metrics=['accuracy'])  
  
# train the model  
# some 10 iterations done here  
history = model2.fit(x_train_scaled, y_train, epochs=10, validation_data=(x_test_scaled, y_test))
```

```

Epoch 1/10
32/32 [=====] - 1s 22ms/step - loss: 1.3086 - accuracy: 0.6980 - val_loss: 0.4044 - val_accuracy: 0.898
9
Epoch 2/10
32/32 [=====] - 0s 14ms/step - loss: 0.1836 - accuracy: 0.9430 - val_loss: 0.3402 - val_accuracy: 0.915
7
Epoch 3/10
32/32 [=====] - 0s 14ms/step - loss: 0.1347 - accuracy: 0.9550 - val_loss: 0.3611 - val_accuracy: 0.882
0
Epoch 4/10
32/32 [=====] - 0s 14ms/step - loss: 0.0675 - accuracy: 0.9840 - val_loss: 0.2386 - val_accuracy: 0.938
2
Epoch 5/10
32/32 [=====] - 0s 14ms/step - loss: 0.0567 - accuracy: 0.9830 - val_loss: 0.2717 - val_accuracy: 0.938
2
Epoch 6/10
32/32 [=====] - 0s 14ms/step - loss: 0.0523 - accuracy: 0.9860 - val_loss: 0.3067 - val_accuracy: 0.932
6
Epoch 7/10
32/32 [=====] - 1s 16ms/step - loss: 0.0499 - accuracy: 0.9860 - val_loss: 0.1749 - val_accuracy: 0.949
4
Epoch 8/10
32/32 [=====] - 0s 15ms/step - loss: 0.0553 - accuracy: 0.9840 - val_loss: 0.3237 - val_accuracy: 0.870
8
Epoch 9/10
32/32 [=====] - 1s 16ms/step - loss: 0.0239 - accuracy: 0.9950 - val_loss: 0.2086 - val_accuracy: 0.943
8
Epoch 10/10
32/32 [=====] - 0s 15ms/step - loss: 0.0129 - accuracy: 0.9990 - val_loss: 0.1806 - val_accuracy: 0.960
7

```

In [117... *# Observation : due to multiple layers the compiling will take more time to execute*
we also got amazing accuracy than earlier

evaluate test dataset on modified model

```
model2.evaluate(x_test_scaled,y_test)
```

```
6/6 [=====] - 0s 5ms/step - loss: 0.1806 - accuracy: 0.9607
```

Out[117]: [0.18058674037456512, 0.9606741666793823]

In [118... *# Earlier we got 0.9213483333587646 now we got 0.9606741666793823 accuracy*


```
# redo the confusion matrix

# build confusion matrix to see how our prediction looks like

# convert to concrete values

y_predicted = model2.predict(x_test_scaled)
y_predicted[0]

y_predicted_labels=[np.argmax(i) for i in y_predicted]

print(y_predicted_labels, len(y_predicted_labels))


conf_mat = tf.math.confusion_matrix(labels=y_test, predictions=y_predicted_labels)
conf_mat
```

```
6/6 [=====] - 0s 5ms/step
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 7, 1, 1, 1, 7, 1, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 9, 4, 2, 2, 5, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
4, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 1, 7, 7, 7, 7, 7,
7, 7, 7, 7, 1, 7, 7, 7, 7, 7, 7, 7, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9,
9, 9, 9, 9, 9, 9, 9] 178
```

```
Out[118]: <tf.Tensor: shape=(10, 10), dtype=int32, numpy=
array([[18,  0,  0,  0,  0,  0,  0,  0,  0,  0],
       [ 0, 14,  0,  0,  0,  0,  0,  2,  0,  0],
       [ 0,  0, 14,  0,  1,  1,  0,  0,  0,  1],
       [ 0,  0,  0, 17,  0,  0,  0,  0,  0,  0],
       [ 0,  0,  0,  0, 18,  0,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0, 18,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0,  0, 18,  0,  0,  0],
       [ 0,  2,  0,  0,  0,  0,  0, 16,  0,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0, 18,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0,  0, 20]])>
```

```
In [119... plt.figure(figsize = (10,10))
sn.heatmap(conf_mat,annot=True,fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
Out[119]: Text(95.72222222222221, 0.5, 'Actual')
```




```
In [106... # Observatoion : we see in the updated model, there are less number of errors,  
# whatever is not in diagonal is a error
```

```
In [120... # Evaluate the model  
test_loss, test_acc = model.evaluate(x_test, y_test)  
print('Test accuracy:', test_acc)
```

```
# Plot the training and validation accuracy  
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('Model accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```

```
6/6 [=====] - 0s 14ms/step - loss: 27.9263 - accuracy: 0.9719  
Test accuracy: 0.9719101190567017
```



In []: